

DraftFM After Release: First Picks in *The Hobbit* and a Preview of *Reality Fracture*

Brian Ward
Independent Researcher

September 25, 2026 — Working draft

Abstract

The original DraftFM paper promised an evaluation of its prerelease *The Hobbit* (HOB) forecast once human drafts became available. This note reports that evaluation. On 6,493 opening packs from experienced winning players, the sealed ranking agrees with 47.34% of first picks, compared with 54.33% for the archived Limited Level-Ups review. DraftFM’s differences from Nizzahon and NicolaiBolas are unresolved by paired intervals; it leads three other archived forecasts. Across all 41,093 shared drafts, DraftFM has the highest agreement among the prerelease comparators. A separate, exploratory model learning just 188 card values from post-release expert picks reaches 68.15% in five-fold held-out evaluation. The results support useful transfer to an unseen set while showing a substantial advantage from observing its drafts. We explain the scope of first-pick evaluation and how a sealed ranking can support actual opening-pack recommendations. We then preview the same frozen model’s *Reality Fracture* (FRA) ranking, including its conspicuous choice of Terminal Criticism as the highest-ranked uncommon. HOB results measure agreement with choices; FRA predictions remain untested.

1 Following through on a prerelease forecast

DraftFM was built for the interval between knowing a new set’s cards and having examples of people drafting them. The original paper describes a model that represents cards through public structured attributes and a fixed text embedding, then scores available choices using the drafted pool and draft context [1]. Its final model was fitted on 32 previously observed expansions, with no HOB draft examples. The HOB card ranking was generated on August 9, 2026 and published under the `draft.fm-v1.0` tag before the August 11 Arena release [2].

That public artifact gives this follow-up a concrete prediction to test. We load the original CSV and verify its digest. We do not rerun the model, choose a newer checkpoint, or revise the ranking using HOB results. The comparison therefore concerns the forecast readers could have obtained before release. The detailed analysis rules were written after release but before calculating the comparison; the forecast is prospective and the analysis is retrospective. This distinction matters more than calling the whole exercise a preregistration.

The main question is modest: when a player opens a fresh pack, how often does the highest-ranked card coincide with the card they take? The answer can be compared with saved human reviews on the same packs. It cannot by itself identify the best possible pick, assign a causal win-rate benefit, or establish how well an assistant manages an entire draft.

2 A common first-pick task

2.1 Why the opening pick is a useful comparison

At pack one, pick one (PIP1), the drafter has no previously selected cards. There is no pool color commitment, curve to repair, or accumulated synergy to favor one card over another. A prerelease card ranking can therefore make a concrete prediction without inventing a way to turn its grades into a full drafting policy. Restrict each ranking to the cards actually offered and take its highest-rated choice.

This also fits the frozen DraftFM architecture. With the empty pool and skill, format, position, and set inputs fixed, each candidate receives its own score. The final export has no whole-set context encoder and no interaction among the offered candidates. Removing alternatives changes the softmax denominator, but does not change the relative order of two remaining cards. Consequently, the published set ranking suffices for this particular first-pick policy. It is not necessary to simulate a draft or rerun the neural network for every opening pack.

First-pick agreement still measures behavior. Strong players can reasonably disagree, select for different strategies, or follow reviews they have seen. The saved reviews and observed players are not necessarily independent sources of judgment. A reviewer who predicts those choices well has passed this benchmark; that does not demonstrate superior game outcomes. Later picks would require pool-conditioned policies for a fair comparison with the full DraftFM model.

2.2 Data, coverage, and ties

The frozen 17Lands Premier Draft file contains 1,833,289 pick rows and 43,650 distinct HOB opening picks, dated August 11–29 [3]. It was downloaded on September 19; the download date is not the end of the observed play period. We use the original paper’s experienced-player thresholds: win-rate bucket at least 0.55 and games-played bucket at least 100. Of the 43,650 drafts, 6,902 satisfy both. This is an operational player slice, not a claim that every selected action is expert ground truth.

Every predictor in a shared comparison must rate every offered non-basic card. Limited Level-Ups marked *The Black Arrow* as sidebar-only, leaving it without a main-deck rank in our transcription. The shared comparison excludes its 2,557 packs, including 409 experienced-player packs. The resulting populations contain 41,093 and 6,493 drafts, respectively. No observed first pick was a basic land or an unknown card. Larger pairwise comparisons remain available in the accompanying results.

The other archived full-set forecasts are NicolaiBolas, Nizzahon, Card Game Base, the Draftsim app pick order, and the Draftsim written review. The two Draftsim forecasts are distinct archived inputs, not two independent organizations. Limited Resources is present in the archive but covers only 120 commons and uncommons. It has no fully covered opening packs and thus no comparable row in the main table.

Letter grades retain their ordinal ordering; numeric scores retain their native values. Slash grades use the first grade under the recorded analysis rule. For offered non-basic card names C_d in draft d , let $T_{m,d}$ be the set tied for the highest rating from predictor m , and let y_d be the observed pick. The per-draft agreement credit is

$$a_{m,d} = \frac{\mathbf{1}\{y_d \in T_{m,d}\}}{|T_{m,d}|}. \quad (1)$$

This is expected agreement under uniform tie-breaking. It avoids favoring a source merely because its scale produces more ties. Card names, rather than duplicate printings of the same card, define the alternatives.

We average credit across the identical eligible drafts for each predictor. Differences receive paired percentile intervals from 2,000 draft resamples, with seed 17. There is exactly one PIP1 row per draft. These are marginal 95% intervals without multiple-comparison adjustment, and draft resampling does not account for repeated appearances of the same player.

3 HOB results: useful transfer, with a clear human lead

Table 1 reports both player populations. The experienced-player slice is primary because it matches the forecast’s intended skill condition. Limited Level-Ups reaches 54.33% agreement, compared with DraftFM’s 47.34%. Its 6.99-point advantage has a paired 95% interval of 5.54–8.38 points: approximately seven additional matching first picks per hundred packs under the tie-breaking convention. That is a meaningful lead on this task.

Predictor	Experienced	All players
HOB-trained choice model, out-of-fold	68.15%	62.01%
Limited Level-Ups (Marc Anderson)	54.33%	49.13%
DraftFM, sealed prerelease	47.34%	50.61%
NicolaiBolas	46.95%	45.78%
Nizzahon	46.67%	47.04%
Card Game Base (Anže Mlakar)	45.56%	43.45%
Draftsim app pick order	45.32%	45.35%
Draftsim review (Andrew Quinn)	43.11%	42.75%
Uniform random	7.42%	7.42%

Table 1: First-pick agreement on identical packs: 6,493 experienced-player drafts and 41,093 all-player drafts. The HOB-trained model uses post-release data with each scored draft held out; it is a separate reference, not a prerelease forecast. Uniform random chooses among distinct offered non-basic card names.

DraftFM’s results against the remaining archived forecasts are more competitive. The paired differences from Nizzahon and NicolaiBolas cross zero (Table 2). That leaves their ordering unresolved at this precision; it does not establish equivalence. DraftFM leads Card Game Base and both Draftsim inputs under the same analysis.

Comparator	DraftFM minus comparator	95% interval
Card Game Base (Anže Mlakar)	+1.79	[+0.51, +3.03]
Draftsim app pick order	+2.02	[+0.82, +3.26]
Draftsim review (Andrew Quinn)	+4.23	[+3.10, +5.43]
Limited Level-Ups (Marc Anderson)	-6.99	[-8.38, -5.54]
NicolaiBolas	+0.39	[-0.84, +1.64]
Nizzahon	+0.67	[-0.35, +1.62]
HOB-trained choice model	-20.81	[-22.09, -19.47]

Table 2: Paired differences on experienced-player packs, in percentage points. Positive values favor DraftFM. Intervals condition on the recorded predictions, including the fitted model’s held-out predictions.

The all-player column reverses the DraftFM–Limited Level-Ups ordering: 50.61% against 49.13%. DraftFM leads all six prerelease comparators there. Reporting both slices reveals this difference rather than selecting the population that gives a preferred story. The present analysis does not isolate why the population changes the ordering.

Coarse review grades deserve a sensitivity check. If any tied-best card instead counted as a full success, Limited Level-Ups would reach 62.88%, while DraftFM would remain at 47.34%. Fractional credit is the headline metric, but the alternative makes its interpretation explicit. We evaluate what the archived grades encode. We cannot recover a creator’s unrecorded preferences among equally graded cards or their complete drafting ability.




Figure 1: Agreement across the two populations. The plotted percentages label the experienced-player slice; paired uncertainty is reported in Table 2.

We also checked secondary outcome associations. The cached card-ratings response contains 188 names but only 18 cards with a non-null game-in-hand win rate and at least 200 observations, and only 52 with usable average-taken-at values and positive pick counts. DraftFM’s Spearman correlations are -0.232 with game-in-hand win rate on the first subset and 0.518 with negative average-taken-at on the second. Those selective samples cannot support a set-wide card-win-rate claim. The complete computed associations are retained in the supplement, including Limited Resources.

4 What observing HOB drafts buys

After the initial comparison, we added a simple set-specific reference. It learns one scalar w_c for each of HOB’s 188 observed non-basic cards. For an offered pack C_d , its choice probabilities are

$$p_w(c \mid C_d) = \frac{\exp(w_c)}{\sum_{j \in C_d} \exp(w_j)}. \quad (2)$$

On training drafts D , it minimizes

$$\sum_{d \in D} \left[\log \sum_{j \in C_d} \exp(w_j) - w_{y_d} \right] + \frac{1}{2} \sum_c w_c^2. \quad (3)$$

These are card identities learned from HOB choices. No card text, structured features, creator grades, game outcomes, or DraftFM warm start enter the fitted values. This is deliberately a small conditional-choice benchmark, not a new neural architecture.

We assign drafts to five folds using CRC32 of the draft ID modulo five. Each fit uses experienced-player drafts from four folds and scores only the other fold. All drafts in a scored fold are excluded from its fit, including when we report the all-player population. The fits use 5,494–5,548 training drafts each. Fixed unit ridge regularization, zero initialization, and L-BFGS-B optimization were specified before computing this addendum’s results; no hyperparameter search was performed. All five fits converged, in 42–45 iterations.

The HOB-trained ranking reaches 68.15% on the 6,493 shared experienced-player packs. Its advantage is 20.81 points over DraftFM (95% interval 19.47–22.09), and 13.82 points over Limited Level-Ups (12.72–14.89). Training agreement ranges from 68.37% to 69.27%; those are diagnostics, not the headline result. The headline predictions are all held out.

The gain shows that post-release choices contain useful information missing from these prerelease rankings, even for a model with only 188 values. It does not show that 188 parameters are generally better than DraftFM’s roughly 1.6 million. The models have different information: the fitted reference has seen thousands of decisions from the target set. Nor is its result a theoretical ceiling. More expressive approaches could perform differently, and matching an observed choice remains distinct from improving a player’s results.

This addendum is exploratory: its recipe followed inspection of the first creator comparison. Random folds mix dates, so it estimates performance on other drafts from this release period rather than future deployment after a particular training cutoff. Its intervals condition on fitted predictions and omit training variance and dependence from overlapping training folds. These qualifications leave the large observed gap intact while defining what the experiment actually tests.

5 Reality Fracture: put another prediction on the table

The next set offers a chance to repeat the exercise. Wizards lists FRA’s Arena release as September 29, with tabletop prerelease beginning September 25 [4]. Using the same frozen checkpoint and feature manifest,

we scored a September 19 card snapshot on September 20. The forecast contains 295 unique names: 285 FRA cards, including five display-only basics, and the ten associated Special Guests. The draftable preview includes FRA and its Special Guests, not the separate Commander product. Alternate printings collapse to one name, using the main set and lowest numeric collector number to select the canonical printing’s features.

Table 3 gives the top ten. Six are planeswalkers, headed by Garruk, Veiled Butcher; the first five appear in Figure 2. Chandra, Torch of Defiance is a main-set reprint in the snapshot (FRA 244), rather than a stray Special Guest. These names make for a lively preview: there is plenty at the top that looks capable of dominating a game. That is a qualitative impression, not a measured forecast of format speed or bomb density. Our letter grades allocate fixed percentile bands, so the number of A’s cannot demonstrate that FRA has more bombs than another expansion.

Rank	Card	Rarity	Grade
1	Garruk, Veiled Butcher	Mythic	A+
2	Ajani Unrelenting	Mythic	A+
3	Garruk, Curse Breaker	Mythic	A+
4	Ajani Resolute	Mythic	A+
5	Sphinx of False Conclusions	Rare	A+
6	Craterclaw Colossus	Mythic	A+
7	The Theorist, Jace Beleren	Mythic	A
8	Uldaros Theorix	Mythic	A
9	Chandra, Torch of Defiance	Mythic	A
10	Guiding Hydra	Rare	A

Table 3: The frozen model’s FRA top ten. Grades are set-relative presentation bands; the underlying raw scores determine the order. The untested forecast was published on September 25, after prerelease began.

5.1 A bold “mythic uncommon” pick

The more entertaining prediction is the uncommon ranking. Among all 109 FRA uncommons, DraftFM puts *Terminal Criticism* first (Table 4; Figure 2). It is a two-mana black instant whose removal is restricted to blue or red creatures and planeswalkers, with one life gained [5]. That is a conspicuously narrow card to nominate as the set’s “mythic uncommon.” The phrase here names the highest-ranked uncommon in the forecast, not a new rarity or a claim about measured wins.

There is a fun connection: six of the model’s own top ten cards are blue or red creatures or planeswalkers that the spell can target. This observation was made after inspecting the ranking. It is not evidence that the network recognized those specific threats and promoted their answer. The frozen export has no whole-set context encoder; its score does not explicitly count those targets or model how often opposing decks will contain them. A reader can reasonably suspect that the model is overvaluing cheap removal and underpricing its restriction. That is precisely why this is a useful prediction to preserve and revisit.

The preview preserves the original prediction instead of editing the bold pick into something more conventional. Raw scores are logits, not card win rates, and gaps between ranks do not supply uncertainty intervals. One feature limitation is also recorded: the frozen structured keyword vocabulary does not include Replicate, although card text still enters through the fixed text encoder. The model and feature manifest have not been retrained on FRA. The exact CSV and a first-pick evaluation recipe were published on September 25 at <https://github.com/brianward92/draftfm/releases/tag/fra-p1p1-2026-09-25>. This is a public timestamp before Arena’s September 29 release, but after prerelease play began. It cannot serve as a pre-prerelease seal.



1. Garruk, Veiled Butcher
A+



2. Ajani Unrelenting
A+



3. Garruk, Curse Breaker
A+



4. Ajani Resolute
A+



5. Sphinx of False Conclusions
A+



Uncommon #1: Terminal Criticism
B+

Figure 2: The five highest-ranked FRA cards and the highest-ranked uncommon in the frozen forecast. Numbers and grades come from the September 20 candidate published on September 25; these are predictions, not measured card outcomes. Images show the canonical FRA printings in the September 19 Scryfall snapshot. Card images are copyright Wizards of the Coast and reproduced at reduced size for scholarly commentary; image source and hashes are recorded with the figure assets.

U rank	Card	Overall rank	Grade
1	Terminal Criticism	49	B+
2	Twisted Fates	54	B+
3	Fblthp, Knows the Way	55	B
4	Jiang Yanggu, Never Alone	57	B
5	Multiply by Zero	61	B
6	Stingerquill Charm	65	B
7	Craftwork Crusher	66	B
8	Essence Burn	67	B
9	Pia, Aether Ascetic	69	B
10	Warrior's Blades	71	B

Table 4: The top ten of all 109 FRA uncommons. This is a filter of the original forecast, with no rescoreing or regrading. The full uncommon ordering and the 295-card forecast accompany the note.

6 From a forecast to a first-pick recommendation

The evaluated policy can be served as a simple lookup. A player supplies an opening pack, through card-name entry or resolved Arena card IDs. Match all offered non-basic cards to the frozen forecast, sort them by raw score, and present the leading choice with two or three alternatives. Show the forecast date and make incomplete card coverage visible. If an offered card has no score, a complete-pack recommendation is unavailable; silently dropping it would change the choice problem.

This workflow needs no new training and no runtime neural inference. Using the archived CSV also preserves the exact rounded scores used in the HOB comparison. It is a practical first release for a manual pack tool and a reference mode for an overlay. Later picks need the drafted pool and the full contextual scorer; a static first-pick list does not provide that context.

For presentation, rank and alternatives are sufficient. A normalized softmax over offered raw scores can describe the model's distribution of choices, but the present evaluation does not establish calibration of those probabilities on an unseen set. A high percentage should not be labeled the probability that a pick is correct, or that it will win. Rules text can explain what a card does; it cannot by itself explain the neural network's reasoning. A close score gap is useful to display as a close comparison, without inventing statistical confidence.

For an application integrating these forecasts, a small versioned data package is enough: card identities, raw scores, ranks, display grades, and the forecast manifest. A reference pack-ranking function can verify that its ordering matches the reported policy. The HOB evaluation and FRA preview then provide both evidence for that policy and a concrete set of predictions to display. Integration should preserve the scored card universe and conditioning settings when recomputing scores; using the same weights alone does not establish parity with an archived forecast.

7 The next experiment

Two follow-ups are especially concrete. First, preserve the published FRA ranking and coverage rules, then freeze creator snapshots and the primary experienced-player P1P1 metric before Arena outcomes arrive. Retain the all-player view and report missing coverage. Preserve Terminal Criticism as the original uncommon nomination; any later diagnosis should explain that prediction rather than replace it.

Second, measure how much set-specific data is needed to improve a frozen ranking. Compare the zero-data forecast, the identity-only choice model, and a model learning corrections to the frozen scores at fixed training sizes, for example 500, 1,000, and 5,000 expert drafts. Use a chronological held-out period and

repeat across sets, with model selection confined to training data. That would test whether transfer reduces the data needed after release. It is a stronger basis for a broader study than the single HOB fitted-model result alone. These are proposed experiments, not results of this note.

The present contribution is the completed forecast check: a fixed prerelease model was competitive with several saved reviews, a strong human review did better on experienced players, and target-set observations improved agreement substantially. The next forecast makes those limits testable again, with a few card predictions worth arguing about along the way.

A Artifacts and reproduction

The accompanying repository contains the analysis protocols, aggregate tables, paired intervals, input hashes, and figure. The HOB forecast CSV has SHA-256

a16c34767aa382e7703f480207f8480484e29ad22bbab2142aa2195e9b9a94a9.

The FRA candidate CSV has SHA-256

116fc528acd83843e283daa73492fa01041f167daebe92103d9f9a52617a348e.

HOB’s raw snapshot and creator input hashes are recorded in `docs/results/hob_oos_20260920/forecast_summary.json`; fitted-model details are in the adjacent `fitted_summary.json`. The full local analysis note, `docs/hob_out_of_sample_20260920.md`, provides executable reproduction commands. The raw draft file is the frozen source of record; later downloads need not be identical.

Independent checks reproduced all 43,650 forecast evaluations and all held-out fitted predictions from saved weights, and verified fold exclusions. The analysis tests cover tie handling, whole-pack coverage, shared populations, name aliases, duplicate drafts, forecast hash checking, the fitted objective’s gradient, and training-fold exclusion. The paper tables are generated from the saved aggregates, with source digests recorded beside them; building the PDF does not refit a model.

Creator grades are the archived pre-Arena transcriptions described in the original paper. Some timestamp evidence is weaker than a verified original publication date: Draftsim’s app date comes from an asset header, and NicolaiBolas’s record establishes retrieval. This supplement reports aggregate comparisons and does not reproduce the creators’ complete grade tables. The included FRA grades are DraftFM’s own outputs. Data attribution: 17Lands.com, CC BY 4.0; card identities and text from Scryfall.

References

- [1] Brian Ward. *DraftFM: A Foundation Model for Day-Zero Drafting in Magic: The Gathering*. August 10, 2026. Original manuscript: `paper/draftfm.tex` in <https://github.com/brianward92/draftfm>.
- [2] Brian Ward. *DraftFM v1.0 sealed HOB forecast*. August 9, 2026. <https://github.com/brianward92/draftfm/tree/draftfm-v1.0>.
- [3] 17Lands. *Public datasets*. HOB Premier Draft snapshot retrieved September 19, 2026. https://www.17lands.com/public_datasets.
- [4] Mike Turian and Jubilee Finnegan. *Collecting Reality Fracture*. Wizards of the Coast, September 8, 2026. Accessed September 20. <https://magic.wizards.com/en/news/feature/collecting-reality-fracture>.
- [5] Scryfall. *Terminal Criticism*, FRA 68. Card record in the September 19, 2026 archived snapshot. <https://scryfall.com/card/fra/68/terminal-criticism>.